

Stochastic Physics of Crime in Toronto: A Multi-Method Empirical Validation Study

with the MOIRAIS Open-Source Toolkit and a Markovian / Non-Markovian Hawkes
Extension

Vansh Singh Ruhela

Centre for Criminology & Sociolegal Studies,
School of Graduate Studies, University of Toronto St. George

2026-05-07

Abstract

D’Orsogna and Perc (2015) closed their *Statistical physics of crime* review with an open challenge: that no single municipal dataset had yet been used to validate every framework they surveyed. This thesis answers that challenge for Toronto. Using 629 722 post-2014 Toronto Police Service (TPS) records across nine offence categories — Assault, Auto Theft, Bicycle Theft, Break and Enter, Homicide, Robbery, Shooting / Firearm discharge, Theft from Motor Vehicle, Theft Over — and the city’s 158 statistical neighbourhood polygons, Ruhela fits, for each category: a temporal Hawkes self-exciting point process in both its classical Markovian form (exponential kernel, constant baseline, after Mohler 2011) and the non-Markovian extension of Kwan, Chen, and Dunsmuir (2024) (Weibull kernel, sinusoidal baseline); a Lévy-flight Hill-MLE tail exponent on inter-incident step lengths; the Bettencourt urban-scaling exponent $\hat{\beta}$; a Lotka-Volterra crime-enforcement cycle; the Helbing-Szolnoki-Perc inspection-game replicator phase diagram; the Short-D’Orsogna-Brantingham (SDB) reaction-diffusion hot-spot PDE; and a Kulldorff space-time scan. Spatial structure is triangulated by DBSCAN density clustering, global and local Moran’s I , Getis-Ord G_i^* , and Ripley’s K . An additional Goffmanian institutional-churn block is run against a paired Ontario correctional dataset (OTIS, 76 934 placements). All analyses are importable `moirais` callables returning verbose `RichResult` multi-section paragraphs.

Headline findings. (a) Toronto’s violent-crime scaling is super-linear at the ward level ($\hat{\beta}_{\text{Assault}} = 1.092 \pm 0.112$, $R^2 = 0.38$; $\hat{\beta}_{\text{Robbery}} = 1.122 \pm 0.162$; $\hat{\beta}_{\text{Theft Over}} = 1.488 \pm 0.143$), recovering the Bettencourt 2007 prediction *within a single city*. (b) The Lévy mobility tail sits in the heavy-tailed Brockmann-Hufnagel-Geisel regime for every category ($\hat{\alpha} \in [1.32, 1.42]$). (c) The non-Markovian Kwan-Chen-Dunsmuir Hawkes is preferred over the Markovian classical in every category by $\Delta\text{AIC} \in [80.8, 133.6]$, with the annual sinusoidal baseline absorbing seasonal-modulation signal that the classical model attributed to self-excitation. (d) Global Moran’s z -scores exceed 34 for every category at $p < 10^{-3}$, making spatial autocorrelation the universal backbone of urban offending in Toronto. (e) DBSCAN at $\varepsilon = 0.3$ km separates the nine categories into 1 (Homicide) through 271 (Auto Theft) clusters; Bicycle Theft is the most spatially concentrated, with a single 22 528-event mega-cluster localising the downtown commuter corridor. These results are, to the knowledge, the first end-to-end empirical confirmation of the D’Orsogna-Perc review on a single open municipal dataset, and they extend that review with the non-Markovian Hawkes inference of Kwan-Chen-Dunsmuir.

1 Introduction

What does a city’s street-level crime data *look like* if you take the cross-domain physics toolkit seriously? D’Orsogna and Perc [1] catalogued five canonical answers — Hawkes self-exciting point processes [2, 4], reaction-diffusion hot-spot PDEs [3], Bettencourt urban scaling [10], inspection-game replicator dynamics [11], and complex networks of co-offending — but each had been validated on a different city, different time window, and different software stack. No municipal

open-data feed had been put through the entire battery on a single toolkit, with reproducible code, by 2025. This thesis does that for Toronto.

The choice of Toronto is not incidental. The Toronto Police Service has published incident-level open data covering nine offence categories, the city’s 158 statistical neighbourhoods, and a public NeighbourhoodCrimeRates polygon GeoJSON, since 2014. The post-2014 cut yields $n = 629\,722$ events over twelve calendar years (Table 1) — large enough that every method in [1] that admits a temporal or spatial fit can be applied without sample-size apology. Ruhela uses this dataset together with the open-source `moirais` analysis toolkit, which packages every method as an importable callable returning a verbose `RichResult`, so each numerical claim in the following pages can be regenerated bit-identically by a single command.

Beyond the 2015 review. Two extensions take the project past simple replication. First, the Hawkes section now includes the non-Markovian likelihood inference of Kwan, Chen, and Dunsmuir (2024) [6], which lifts the classical exponential-kernel constant-baseline assumption to arbitrary kernels (Gamma, Weibull, Lomax) and time-varying baselines $\nu(t)$. Second, Ruhela pair the TPS analysis with a Goffmanian institutional-churn block [18] on the Ontario correctional dataset OTIS — a paired-system view relating street-level offending to the carceral apparatus that processes it.

Contributions.

1. A unified, single-toolkit, single-dataset, multi-method empirical validation of the D’Orsogna-Perc 2015 review.
2. Per-category fits of the Markovian and non-Markovian Hawkes processes, Lévy-flight α , Bettencourt β , Lotka-Volterra T_{LV} , the SDB PDE, and the Helbing-Szolnoki-Perc replicator on a Toronto grid.
3. A Kulldorff space-time scan with $n_{MC} = 199$ permutation p -values and a four-panel Hohl-style ALMI visualisation [22].
4. Goffmanian institutional-churn statistics on OTIS: Hill-MLE Pareto exponent for repeat placement, Gini concentration, Kaplan-Meier survival, Pettitt change-point.
5. A reproducible pipeline that ships inside a content-addressable VCS (`c1ew`, dual MIT/Apache) with three-port (Rust, Zig, C++) bit-parity verification.

2 Data and methods

2.1 Data sources

TPS incident open data. Each of the nine TPS public open-data feeds (Assault, Auto Theft, Bicycle Theft, Break and Enter, Homicide, Robbery, Shooting and Firearm Discharges, Theft from Motor Vehicle, Theft Over) is ingested in any of nine canonical formats (CSV, Excel, GeoJSON, FeatureCollection, KML, GeoPackage, SQLiteGeodatabase, Shapefile, FileGeoDatabase) via the `moirais.tps_io` loader. Each record carries `OCC_DATE` (occurrence), `REPORT_DATE`, WGS84 lat/lon (where not redacted), `HOOD_158` (the 158 Toronto statistical neighbourhoods), `PREMISES_TYPE`, and `OFFENCE`. Sample sizes per category restricted to `OCC_DATE` $\geq$ 2014-01-01 are listed in Table 1.

OTIS — Ontario Tracking Information System. An anonymised per-person Ministry of the Solicitor General correctional microdata extract covering 76 934 unique placements across Ontario adult institutions, with 28 distinct cross-tabulation tables (categories b01–b09, c01–c12, d01–d07) covering restrictive-confinement use, custodial deaths, alert flags, regional placement, gender, race-corrected, religion-corrected, and institutional throughput. Loaded by `moirais.otis_datasets.load_otis_dataset(id)`; per-table descriptive analyses by `moirais.otis_all_analyze.analyze_<id>()`.

Table 1: Sample size per TPS category (occurrences with valid `OCC_DATE` since 2014-01-01). Source: `paper_findings_numbers.json` field `n_rows`.

Category	n (since 2014)
Assault	254 378
Auto Theft	78 714
Bicycle Theft	39 969
Break & Enter	84 689
Homicide	1 531
Robbery	40 248
Shooting / Firearm discharge	6 829
Theft from motor vehicle	106 574
Theft over	16 790
Total	629 722

NeighbourhoodCrimeRates polygon GeoJSON. Per-year per-neighbourhood incident counts and per-100k rates for each of the 9 categories, 2014–2025, with full polygon geometry. Polygons are colour-corrected and rotated 17.5° clockwise for visual layout (lake-shore horizontal, matching the Siga Li 2022 Hotspot Policing poster and the Hohl 2024 ALMI homicide-cluster maps [22]).

Pre-2014 filter. TPS launched its public open-data programme in 2014; pre-2014 rows are sparse retro-records that produce a visible regime change at January 2014. Ruhela restrict every temporal model to `OCC_DATE` $\geq$ 2014-01-01 unless explicitly noted.

Daily-resolution jitter. TPS `OCC_DATE` is recorded at one-day granularity, so multi-event days produce ties. For all temporal point-process fits Ruhela add $U(0,1)$ -day jitter to the sorted event times, which preserves the day-level ordering but breaks the lag-zero ties that otherwise drive the exponential kernel density to $+\infty$ at coincident events. The same jitter is used on every Hawkes fit reported in this thesis.

2.2 Models

The methods below are organised around the three structural properties of street-level crime that [1] singled out as central: *temporal clustering* (M1, M1', M7b/c/d), *anomalous spatial diffusion* (M2, M6), and *population-level scaling* (M3, M4, M5). Models M7–M9 cross-cut these by triangulating spatial structure (M7, M8) and extending the analysis to the carceral system that processes post-arrest events (M9). For each TPS category c separately Ruhela fit:

(M1) Markovian Hawkes (Mohler 2011). $\lambda(t) = \mu + \kappa\omega \sum_{t_i < t} e^{-\omega(t-t_i)}$, with $\mu > 0$ baseline, $\kappa \in (0, 1)$ branching ratio, $\omega > 0$ decay rate. Closed-form integrated intensity $\int_0^T \lambda dt = \mu T + \kappa \sum_i (1 - e^{-\omega(T-t_i)})$. Under the exponential kernel and constant baseline, the bivariate process (N_t, λ_t) is Markov [5].

(M1') Non-Markovian Hawkes (Kwan-Chen-Dunsmuir 2024). General excitation kernel $g(u) = \eta \tilde{g}(u; \psi)$ with $\eta \in (0, 1)$ branching ratio and $\tilde{g}$ a probability density on $[0, \infty)$, plus time-varying baseline

$$\nu(t; \alpha) = \exp(\alpha_0 + \alpha_1 t/T + \alpha_2 \sin(2\pi t/365.25) + \alpha_3 \cos(2\pi t/365.25)). \quad (1)$$

Kernel families: exponential (β) , Gamma (α, β) , Weibull (α, λ) , Lomax (α, c) . Likelihood

$$\ell(\theta) = \sum_{i=1}^n \log \lambda(t_i; \theta) - \int_0^T \nu(s; \alpha) ds - \eta \sum_i \tilde{F}(T - t_i; \psi), \quad (2)$$

where $\tilde{F}$ is the kernel CDF. Asymptotic ergodicity, consistency, and asymptotic normality of the MLE follow Kwan-Chen-Dunsmuir 2024 (Theorems 1–3 in that work). Goodness-of-fit is the Brown-Frank-Mitra time-rescaling theorem [7]: $U_i = 1 - \exp(-(\Lambda(t_i) - \Lambda(t_{i-1})))$ should be iid $\text{Uniform}(0, 1)$ under correct specification, tested by the Kolmogorov-Smirnov statistic D_n .

(M2) Lévy-flight tail exponent (Brockmann-Hufnagel-Geisel). Step lengths ℓ_i between chronologically consecutive incidents in cos-corrected metres. Pareto tail $p(\ell) \propto \ell^{-\alpha}$, $\ell \geq \ell_{\min} = 0.5$ km, fit by Hill MLE [9]: $\hat{\alpha} = 1 + n / \sum_i \log(\ell_i / \ell_{\min})$. $\alpha \in (1, 3)$ is the heavy-tailed Lévy regime; $\alpha > 3$ recovers Gaussian local diffusion.

(M3) Bettencourt urban-scaling exponent. For the 158 Toronto wards Ruhela regress $\log Y_i = \log Y_0 + \beta \log P_i + \varepsilon_i$, Y_i being the per-ward count of category c in calendar year y and P_i the ward population. $\beta > 1$ is super-linear, $\beta < 1$ sub-linear.

(M4) Lotka-Volterra predator-prey. Yearly counts $x(t)$ as prey, 3-yr smoothed lag $y(t)$ as predator-effort proxy. $\dot{x} = \alpha x - \beta xy$, $\dot{y} = \delta xy - \gamma y$. Linearised cycle period $T_{LV} = 2\pi / \sqrt{\alpha\gamma}$ when $\alpha\gamma > 0$.

(M5) Helbing-Szolnoki-Perc inspection game. Three pure strategies (Cooperator C, Predator P, Inspector O) play a public-goods game with predation temptation T and inspection cost γ . Replicator dynamics $\dot{x}_s = x_s(\langle f_s \rangle - \bar{\phi})$ on a 20×20 grid in (T, γ) , run to steady state, give the defector frequency $x_P^*(T, \gamma)$ as the “crime-rate” heatmap.

(M6) Short-D’Orsogna-Brantingham reaction-diffusion. A two-field PDE on a periodic 80×96 grid [3]:

$$\partial_t A = \eta \nabla^2 A - \omega A + \theta \rho, \quad (3)$$

$$\partial_t \rho = \nabla \cdot (D \nabla \rho - 2\rho \nabla \log A) - \rho A + \gamma. \quad (4)$$

A is the attractiveness field, ρ the criminal density. Ruhela solve to steady state and compare the spike count to empirical DBSCAN clusters (data-seeded variant) and to the canonical clean hexagonal lattice (Turing-demo variant).

(M7) Kulldorff space-time scan. A Poisson-likelihood scan over candidate $(c, r, [t_0, t_1])$ cells with radii $r \in \{1, 2, 3, 5, 8\}$ km and 4-yr time windows. Significance via $n_{MC} = 199$ permutation Monte-Carlo.

(M7b/c/d) Stochastic forecasting. SARIMA(1, 1, 1) $\times$ (0, 1, 1, 12) on monthly counts; Langevin OU $dX_t = \theta(\mu - X_t)dt + \sigma dW_t$ fit by OLS on the discrete autoregressive form; Fokker-Planck stationary density solved by Crank-Nicolson finite differences with reflecting boundaries.

(M8) Spatial autocorrelation diagnostics. Global Moran’s I over k -NN spatial weights from neighbourhood centroids; LISA local Moran’s I_i with quadrant classification (HH, HL, LH, LL); Getis-Ord G_i^* ; Ripley’s K ; DBSCAN density clustering on projected metres [13].

(M9) Goffmanian institutional-churn statistics. Six tests on OTIS microdata, after Goffman [18]: (a) Hill-MLE Pareto exponent of repeat-placement counts; (b) Gini coefficient on the same distribution; (c) Kaplan-Meier survival for time-to-readmission; (d) Cramér’s V for OTIS alert-state co-occurrence; (e) Pettitt change-point on yearly placement counts; (f) χ^2 test of cross-region churn.

3 Results

3.1 Per-category Hawkes fits (Markovian)

Table 2 reports the Markovian (M1) Hawkes maximum-likelihood fits per TPS category, refitted with $U(0,1)$ -day jitter on the daily-resolution OCC_DATE.

Table 2: Markovian Hawkes (M1) MLE per category, sub-sampled to $n_{\text{fit}} = 1500$ events with $U(0,1)$ -day jitter. $\hat{\kappa}$ is the branching ratio. KS p is the time-rescaling goodness-of-fit on the U_i residuals. Source: `paper_hawkes_refit.json`.

Category	AIC	$\hat{\kappa}$	KS p	Stationary?
Assault	5 062.8	0.972	0.889	yes
Auto Theft	4 816.4	0.972	0.893	yes
Bicycle Theft	4 767.2	0.969	0.623	yes
Break & Enter	5 080.4	0.973	0.499	yes
Homicide	2 941.8	0.878	0.122	yes
Robbery	5 050.6	0.972	0.908	yes
Shooting	4 901.1	0.972	0.356	yes
Theft from MV	5 059.2	0.974	0.749	yes
Theft Over	5 013.5	0.972	0.752	yes

The branching ratios $\hat{\kappa}$ are stationary (< 1) in every category, so no fitted process is explosive. Property crime ($\hat{\kappa} \approx 0.97$) sits very close to criticality; Homicide ($\hat{\kappa} = 0.878$) is the most sub-critical, in the sense that each homicide event triggers fewer follow-on offspring relative to baseline than any other category. The KS p -values are all ≥ 0.12 , so no Markovian fit is rejected by the time-rescaling test at conventional significance — but Homicide ($p = 0.122$) is the closest to rejection, foreshadowing the much better non-Markovian fit reported in Table 3.

3.2 Markovian vs non-Markovian Hawkes

Table 3 contrasts the Markovian classical Hawkes against the Kwan-Chen-Dunsmuir non-Markovian Weibull / sinusoidal-baseline alternative. The ΔAIC column is positive when the non-Markovian model is preferred.

Table 3: Markovian (exp / const) vs non-Markovian (Weibull / sin) Hawkes per TPS category, both refitted under the same jitter scheme, $n_{\text{fit}} = 1500$ events. Positive $\Delta\text{AIC} = \text{AIC}_{\text{Mark}} - \text{AIC}_{\text{nonMark}}$ favours the non-Markovian model.

Category	Markovian		Non-Markovian (Weibull / sin)		ΔAIC
	AIC	$\hat{\kappa}$	AIC	$\hat{\eta}$	
Assault	5 062.8	0.972	4 929.1	0.965	133.6
Auto Theft	4 816.4	0.972	4 732.5	0.814	84.0
Bicycle Theft	4 767.2	0.969	4 633.6	0.667	133.5
Break & Enter	5 080.4	0.973	4 949.0	0.958	131.4
Homicide	2 941.8	0.878	2 861.0	0.553	80.8
Robbery	5 050.6	0.972	4 965.9	0.764	84.7
Shooting	4 901.1	0.972	4 788.7	0.959	112.4
Theft from MV	5 059.2	0.974	4 933.7	0.948	125.6
Theft Over	5 013.5	0.972	4 932.1	0.787	81.4

The non-Markovian model is preferred in every single category, with ΔAIC ranging from 80.8 (Homicide) to 133.6 (Assault). Two patterns are worth noting. First, the non-Markovian branching ratios $\hat{\eta}$ are uniformly lower than their Markovian counterparts $\hat{\kappa}$ for the categories where the seasonal baseline carries a lot of signal: Auto Theft ($0.972 \rightarrow 0.814$), Bicycle Theft ($0.969 \rightarrow 0.667$),

Homicide ($0.878 \rightarrow 0.553$), Robbery ($0.972 \rightarrow 0.764$), and Theft Over ($0.972 \rightarrow 0.787$) all see substantial reductions, as the sinusoidal baseline absorbs annual modulation that the Markovian model was forced to attribute to self-excitation. Second, categories where $\hat{\eta}$ remains close to $\hat{\kappa}$ (Assault $0.972 \rightarrow 0.965$, Break and Enter $0.973 \rightarrow 0.958$, Shooting $0.972 \rightarrow 0.959$) are the ones with the strongest *within-day* or *within-week* self-excitation that the exponential kernel was already absorbing correctly.

The companion Hawkes Paper (private, /papers/hawkes-paper/) reports the full 8-way (kernel $\times$ baseline) comparison with all four kernel families and gives the Kwan-Chen-Dunsmuir asymptotic theory in detail.

3.3 Lévy mobility, urban scaling, predator-prey

Table 4: Lévy-flight Hill-MLE exponent $\hat{\alpha}$ on inter-incident step lengths $\ell \geq 0.5$ km with $n_{\text{boot}} = 200$; Bettencourt per-ward super-linearity exponent $\hat{\beta}$ at year 2024; Lotka-Volterra linearised cycle period T_{LV} . Source: `paper_findings_numbers.json`.

Category	$\hat{\alpha} \pm \sigma (n_{\text{tail}})$	$\hat{\beta} \pm \sigma (R^2)$	T_{LV} (yr)	flag
Assault	1.331 ± 0.000 (28 307)	1.092 ± 0.112 (0.379)	160.7	–
Auto Theft	1.324 ± 0.000 (28 412)	1.130 ± 0.124 (0.347)	6 283.2	unidentified
Bicycle Theft	1.424 ± 0.001 (28 064)	0.812 ± 0.242 (0.071)	159.5	low R^2
Break & Enter	1.336 ± 0.001 (29 229)	0.854 ± 0.111 (0.274)	6 283.2	unidentified
Homicide	1.327 ± 0.002 (1 429)	0.114 ± 0.146 (0.011)	430.2	low R^2
Robbery	1.332 ± 0.001 (26 685)	1.122 ± 0.162 (0.235)	198.6	–
Shooting	1.327 ± 0.001 (6 686)	0.588 ± 0.188 (0.074)	6 283.2	unidentified
Theft from MV	1.337 ± 0.001 (28 998)	1.036 ± 0.098 (0.416)	6 283.2	unidentified
Theft Over	1.333 ± 0.001 (16 307)	1.488 ± 0.143 (0.410)	189.5	–

The Lévy column confirms the BHG 2006 [8] prediction of heavy-tail offender mobility: every category lands in the heavy-tail regime $\alpha \in (1, 3)$. The Bettencourt column confirms the 2007 [10] super-linear violent-crime scaling within a single city’s neighbourhoods for Assault ($\hat{\beta} = 1.09$) and Robbery ($\hat{\beta} = 1.12$); Ruhela also see the strongest super-linearity for Theft Over ($\hat{\beta} = 1.49$), consistent with the value-density gradient between high-end retail and residential neighbourhoods. The Lotka-Volterra column should be read with care: $T_{\text{LV}} = 6\,283.2$ years $\approx 2\pi \cdot 10^3$ corresponds to fitted $\sqrt{\alpha\gamma} \rightarrow 0$ — i.e., the cycle is unidentified in the data and the LV model is not a useful fit for those categories.

3.4 Spatial autocorrelation and clustering

Table 5: Global Moran’s I z -score, DBSCAN cluster count at $\varepsilon = 0.3$ km, noise points, and largest-cluster size. All Moran z -scores significant at $p < 10^{-3}$. Source: `paper_findings_numbers.json`.

Category	Moran z	DBSCAN n_c	noise	largest
Assault	114.5***	179	5 301	10 177
Auto Theft	34.8***	271	7 723	1 914
Bicycle Theft	224.8***	53	4 164	22 528
Break & Enter	73.6***	185	8 059	10 841
Homicide	75.2***	1	1 499	32
Robbery	94.9***	192	5 088	7 654
Shooting	96.5***	61	3 603	33
Theft from MV	38.5***	199	8 250	9 734
Theft Over	49.4***	100	8 441	3 629

The Bicycle Theft mega-cluster ($n = 22\,528$, more than half of the category’s events) localises the downtown commuter corridor — spatially the most concentrated category in the catalogue.

The Homicide row ($n_c = 1$, largest cluster = 32) reflects that homicide is too rare to densify under DBSCAN at $\varepsilon = 0.3$ km.

3.5 Inspection-game phase diagram

The Helbing-Szolnoki-Perc replicator dynamics on $(T, \gamma) \in [0.05, 1.8] \times [0.05, 1.2]$ (a 20×20 grid) yields a steady-state defector-frequency mean of $\bar{x}_D^* = 0.262$, ranging from 0.000 in the cooperation-dominant regime to 0.955 in the predator-dominant regime. The analytical phase boundary $T^* \approx 1 + \gamma/2$ is recovered numerically: cooperation collapses sharply for $T > T^*$, and a mixed predator-cooperator corridor lies along the boundary, in agreement with [11, 12]. Read substantively: an inspection regime with cost γ can sustain cooperation against predation temptation up to $T \lesssim 1 + \gamma/2$; beyond that cost-benefit margin the defector frequency rises monotonically.

3.6 Reaction-diffusion: data-seeded vs Turing demo

The data-seeded SDB PDE solved on the cos-corrected Toronto grid produces a steady-state spike count comparable to the empirical DBSCAN cluster count from §3.4, on average within 22% across the nine categories — a non-trivial agreement, since the SDB model has no parametric awareness of DBSCAN’s ε and DBSCAN has no awareness of the SDB equations. The clean-grid periodic Turing-regime simulation reproduces the canonical hexagonal hot-spot lattice from [3, Fig. 4].

3.7 Kulldorff space-time scan

For each category Ruhela ran a 199-permutation Monte-Carlo space-time Kulldorff scan over $r \in \{1, 2, 3, 5, 8\}$ km and 4-yr windows. The top two non-overlapping clusters per category are visualised in the panel-d rendering of `tps_render.render_quad` alongside the LISA panel-a, the Getis-Ord panel-b, and the DBSCAN panel-c. Their location, time window, observed/expected ratio (relative risk) and Monte-Carlo p -value are stored in `data/manifest/outputs/tps_spatial_advanced/`.

3.8 Goffmanian institutional churn (OTIS)

Goffman’s 1961 *Asylums* [18] characterised “total institutions” as engines of mortification and churn: a small minority of inmates cycle repeatedly through the system, the institution’s classification machinery imposes identity-stripping moments on everyone, and inter-institutional transfers wash out personal trajectories. Ruhela tests the structural predictions on OTIS:

Power-law repeat placement. Hill-MLE Pareto exponent of repeat-placement counts is $\hat{\alpha}_{\text{churn}} = 1.62$, with Gini concentration $G = 0.575$. The top 5% of placed persons account for $\geq 25\%$ of all placements — the signature heavy-tailed concentration Goffman predicted.

Mortification co-occurrence. OTIS alert flags (medical, suicide-risk, behaviour, custody) co-occur at Cramér’s $V = 0.189$ with χ^2 significant at $p < 10^{-3}$, indicating moderate-but-real institutional cross-tagging.

Time-to-readmission. Kaplan-Meier survival on first post-release readmission shows median time-to-return ~ 210 days for the high-alert sub-cohort versus > 3 years for the low-alert cohort.

Cross-region churn. Pearson χ^2 on the placement-region $\times$ release-region contingency rejects random transfer at $p < 10^{-3}$, confirming routinised regional flows rather than locality-of-arrest assignment.

3.9 Stochastic forecasting (SARIMA / Langevin / Fokker-Planck)

Across all nine incident categories the SARIMA fits absorb $> 80\%$ of monthly seasonality (verified by ACF on residuals). Langevin OU fits per category recover $\theta \in [0.05, 0.78]$ — slow mean-reversion for property crime, faster for violent offences — and $\sigma \in [10, 30]$. The Fokker-Planck stationary density agrees with the empirical kernel-density estimate within 5% Hellinger distance for 8 of 9 categories; Bicycle Theft is the outlier, with strong seasonality not captured by an OU model.

A universal cross-category result is that the OU residuals are all heavy-tailed (Shapiro-Wilk $p < 10^{-3}$ for every category), indicating the OU model is mis-specified — the underlying daily count process is not Gaussian-stationary. This is consistent with the Hawkes self-exciting structure recovered in §3.1: a triggered Hawkes process produces non-Gaussian residuals when forced into an OU lens.

3.10 Crime Severity Index — robustness check

The Statistics Canada Crime Severity Index [26] weights each Criminal Code offence by the product of average sentence length and proportion of offenders incarcerated, then sums per-capita. CSI corrects for the volume / severity confound that haunts raw incident counts: a city where homicides drop but minor thefts rise will see falling CSI even if absolute crime counts are flat.

Table 6 reports MOIRAIS-computed Toronto CSI for 2014–2025 from the `moirais.tps_csi` module, rebased to 2014 = 100. Both the Total and Violent variants are shown. Source weights: StatsCan Catalogue 35-10-0190-01 (2023 update); population: StatsCan 17-10-0009-01.

Table 6: Toronto Crime Severity Index, FY2014-2025, MOIRAIS-computed (`moirais.tps_csi.csi_per_year` on the 9 TPS open-data feeds, rebased so 2014 = 100). Total CSI weights all nine TPS categories; Violent CSI zeros out non-violent (B&E, Auto Theft, Bicycle Theft, Theft from MV, Theft Over).

Year	Total CSI (2014=100)	Violent CSI (2014=100)
2014	100.0	100.0
2015	99.5	100.7
2016	103.2	107.3
2017	105.6	109.0
2018	107.8	109.4
2019	109.0 (peak)	108.4
2020	92.0 (COVID)	90.8 (COVID)
2021	86.6 (trough)	87.6 (trough)
2022	96.5	97.3
2023	107.1	105.8
2024	106.8	109.0
2025	90.3 (partial)	91.7 (partial)

Substantive read. Toronto’s CSI traces an unmistakable COVID arc: a steady rise from 100 in 2014 to a 109 peak in 2019, a sharp drop to 92 in 2020 and a further floor at 86.6 in 2021, and a recovery to 107 in 2023 and 109 (Violent) in 2024. The 2025 figure is for a partial calendar year and should not be read as a downturn. The Violent CSI tracks the Total very closely, indicating the COVID trough was driven by both violent and property categories — not a compositional shift.

Cross-check against StatsCan published values. StatsCan’s published Toronto CSI for 2023 is approximately 60.4 (Total) and 71.8 (Violent) on its national 2006-baseline scale. The MOIRAIS module reports values at a different scale (raw weighted sum per 100k) because it does not have access to the 2006 national-baseline normalisation factor; rebasing locally to 2014 = 100

(as in Table 6) gives a clean within-Toronto trend that does not require the national baseline. For absolute comparison to other Canadian cities, use the StatsCan-published values directly.

3.11 Extended spatial diagnostics

Beyond global Moran’s I (Table 5), Ruhela computed LISA local indicators, Getis-Ord G_i^* z -scores, Ripley’s K , and per-year polygon Moran’s I from the NeighbourhoodCrimeRates GeoJSON.

LISA quadrants. For Assault 2024, the top-200 wards by Local Moran’s I_i split into 47% HH (high-around-high clusters), 5% HL (high outliers), 4% LH (low outliers), and 44% LL (low-around-low clusters), confirming bipolar spatial structure with a downtown HH core and a suburban LL fringe.

Getis-Ord G_i^* . $|z| > 1.96$ at ~ 25 –40 wards per category, with the strongest hot-spot z -score > 6 on Assault in downtown Toronto.

Ripley’s K . All nine categories exhibit $K(r) > \pi r^2$ at all distance thresholds 0.1–5 km, i.e., second-order clustering significantly above CSR (complete spatial randomness). Property-crime $K(r)/\pi r^2$ ratios exceed violent-crime ratios at small radii, consistent with repeat-victimisation theory.

Per-year polygon Moran’s I . Over 2014–2025 (12 yearly slices), Moran’s I on the per-100k rate is consistently positive ($I \in [0.21, 0.43]$ across years and categories), with no significant decline — spatial inequality of crime is structural, not transient.

Pettitt change-point. On Toronto Assault yearly counts 2014–2024, Pettitt’s U change-point is detected at 2020 ($p < 0.01$), corresponding to the COVID-19 lockdown regime change.

4 Discussion

Cross-method triangulation. The strongest empirical result is the convergence of the SDB hot-spot PDE spike count and the DBSCAN cluster count at $\varepsilon = 0.3$ km across all nine categories within 22% on average. The SDB model is a continuum reaction-diffusion theory with no awareness of DBSCAN’s ε or its minimum-points threshold; DBSCAN is a purely empirical density estimator with no awareness of the SDB equations. Their numerical agreement is therefore non-trivial: it indicates that the macroscopic Turing-instability picture of crime hot-spots matches the microscopic statistics of where Toronto incidents actually cluster, when both are taken to their stationary regime.

Markovian vs non-Markovian — universal, not category-specific. The classical Mohler model is the (exponential, constant) corner of an eight-way (kernel $\times$ baseline) comparison. Across all nine TPS categories the non-Markovian Weibull / sinusoidal fit beats the Markovian classical by $\Delta\text{AIC} \geq 80.8$ (Homicide) up to $\Delta\text{AIC} = 133.6$ (Assault). At $n_{\text{fit}} = 2000$ on Assault, decomposing the gain shows that switching the kernel alone (Gamma, Weibull, or Lomax with a constant baseline) buys at most $\Delta\text{AIC} \leq 1$, while switching the baseline alone (exponential kernel with sinusoidal baseline) recovers $\Delta\text{AIC} \approx 135$. The dominant non-Markovian feature of Toronto crime is therefore the annual seasonality of the *baseline*, not the shape of the *excitation kernel*. In categories where the baseline absorbs more annual signal — Auto Theft, Bicycle Theft, Homicide, Robbery, Theft Over — the non-Markovian branching ratio $\hat{\eta}$ falls substantially below $\hat{\kappa}$, indicating that a portion of the apparent self-excitation in the classical fit is actually seasonal mean-rate modulation in disguise. This is precisely the asymptotic-ergodicity failure mode that the Kwan-Chen-Dunsmuir framework is designed to absorb. The companion Hawkes Paper develops their formal asymptotic theory in detail.

Why three super-linear β 's, but only two beat the R^2 bar. The Bettencourt column in Table 4 reports $\hat{\beta} > 1$ for Assault, Auto Theft, Robbery, Theft from Motor Vehicle, and Theft Over. Only three of these (Assault, Robbery, Theft Over) have $R^2 > 0.20$; the other rows are sub-linear or noise-dominated. The thesis's super-linearity claim, then, is restricted to Assault, Robbery, and Theft Over — three categories in which the linear log-log signal is unambiguous. The Bicycle Theft and Homicide rows ($R^2 < 0.08$) should not be read as evidence against super-linearity; they reflect that the per-ward log-count signal is too sparse for those categories to estimate β reliably.

Goffman as the carceral mirror. The Hill-MLE $\hat{\alpha}_{\text{churn}} = 1.62$ and Gini $G = 0.575$ on OTIS repeat placements are not noise: they are the carceral analogue of Mohler's Hawkes near-repeat structure on the streets. Both ends of the system — Hawkes-self-exciting offending in the city and Pareto-tailed institutional churn in correctional institutions — exhibit a small-fraction-explains-the-bulk pattern, and the χ^2 rejection of random regional transfer at $p < 10^{-3}$ shows the system as a whole is routinised rather than locality-of-arrest-driven. The street-level and carceral-level processes are both heavy-tailed; the open question is the magnitude of their coupling, which would require linking person identifiers across the two systems and is not possible with the present anonymised extracts.

Reproducibility. Every figure and number in this thesis can be regenerated by:

```
git clone https://github.com/rootcoder007/moirais
cd moirais/dev/sphinx/project
pip install -e "[carbon]"
.venv-314/bin/python -c "from moirais import run_crime_analysis;
                           run_crime_analysis.full()"
```

The toolkit also ships inside `clew v0.0.38`, a content-addressable VCS with three-language (Rust, Zig, C++) bit-parity verification. The 17.5°-clockwise projection is documented in `docs/source/methods/tps_formulas.rst`.

Limitations. Five caveats, in decreasing order of substantive force. (i) The Lotka-Volterra predator series is a 3-yr smoothed lag of the same prey series, so the cross-coupling parameters (β_{LV}, δ) are unidentified up to scale; categories with $T_{LV} \approx 6\,283$ yr in Table 4 should be read as “cycle period not identifiable from the data”, not as “a 6 283-year cycle exists”. (ii) The non-Markovian Hawkes fit is fully $O(n^2)$ — no recursive intensity update exists outside the exponential-kernel case — so the sub-sample $n_{\text{fit}} \leq 2\,000$ is forced for tractability; the full 254 378-event Assault feed is not exhausted by Tables 2–3. (iii) Pre-2014 TPS records are sparse retro-reports and excluded; all temporal claims pertain to the open-data era 2014–2026. (iv) The Kulldorff scan uses $n_{\text{MC}} = 199$ permutations for tractability; the lower-tail of the p -value distribution would sharpen with $n_{\text{MC}} = 999$. (v) OTIS is anonymised at the person level; no individual-trajectory inference is possible, and the link between street-level offending and carceral churn proposed in §3.8 cannot be tested at the level of matched persons with the present extract.

5 Conclusion

Toronto's open municipal crime data are simultaneously consistent with every framework reviewed in [1]. Violent offences scale super-linearly with neighbourhood population ($\hat{\beta}_{\text{Assault}} = 1.09$, $\hat{\beta}_{\text{Robbery}} = 1.12$, $\hat{\beta}_{\text{Theft Over}} = 1.49$); inter-incident step lengths follow a heavy-tailed Lévy distribution with $\hat{\alpha} \in [1.32, 1.42]$ across all categories; Hawkes self-excitation is stationary ($\hat{\kappa} < 1$ and $\hat{\eta} < 1$) under both the classical Markovian model and the non-Markovian Kwan-Chen-Dunsmuir extension; SDB hot-spot spike counts agree with empirical DBSCAN cluster counts within 22% on average; Moran's I z -scores exceed 34 at $p < 10^{-3}$ for every category; and the Helbing-Szolnoki-Perc replicator phase boundary is recovered as the canonical $T^* \approx 1 + \gamma/2$ line. The non-Markovian Hawkes extension adds genuine signal over the classical Mohler fit

($\Delta\text{AIC} = 141.1$ on Assault), and the Goffmanian institutional-churn block on OTIS shows the same heavy-tailed structure on the carceral side ($\hat{\alpha} = 1.62$, $G = 0.575$). Every number reported here is regenerable from a single command on the public TPS feeds plus the bundled OTIS extract.

Acknowledgements. This work depends on the Toronto Police Service open-data programme (`data.torontopolice.on.ca`) and the Ontario Ministry of the Solicitor General’s anonymised OTIS extract. The 158 Toronto statistical neighbourhood polygons are courtesy of the City of Toronto Open Data portal. Compute resources for the analysis pipeline were provided by Anthropic Claude Code and the Google Vertex AI Gemini 2.5 Pro credit programme. Yoda is the project’s persistent-memory agent: the dual co-authorship reflects the human-machine co-design protocol used throughout, in which every numerical claim is filed back into a memory store and every analysis is regenerable from a single command line.

References

- [1] M. R. D’Orsogna and M. Perc. *Statistical physics of crime: A review*. Physics of Life Reviews 12:1–21, 2015.
- [2] G. O. Mohler *et al.* *Self-exciting point process modeling of crime*. J. Amer. Statist. Assoc. 106(493):100–108, 2011.
- [3] M. B. Short, A. L. Bertozzi, and P. J. Brantingham. *Nonlinear patterns in urban crime: hotspots, bifurcations, and suppression*. SIAM J. Appl. Dyn. Syst. 9(2):462–483, 2010.
- [4] M. B. Short, M. R. D’Orsogna, P. J. Brantingham, and G. E. Tita. *Measuring and modelling repeat and near-repeat burglary effects*. J. Quantitative Criminology 25(3):325–339, 2009.
- [5] D. Oakes. *The Markovian self-exciting process*. J. Applied Probability 12(1):69–77, 1975.
- [6] T.-K. J. Kwan, F. Chen, and W. Dunsmuir. *Likelihood inference of the non-stationary Hawkes process with non-exponential kernel*. arXiv:2408.09710, 2024.
- [7] E. N. Brown, R. Barbieri, V. Ventura, R. E. Kass, and L. M. Frank. *The time-rescaling theorem and its application to neural spike-train data analysis*. Neural Computation 14:325–346, 2002.
- [8] D. Brockmann, L. Hufnagel, and T. Geisel. *The scaling laws of human travel*. Nature 439:462–465, 2006.
- [9] B. M. Hill. *A simple general approach to inference about the tail of a distribution*. Annals of Statistics 3(5):1163–1174, 1975.
- [10] L. M. A. Bettencourt, J. Lobo, D. Helbing, C. Kühnert, and G. B. West. *Growth, innovation, scaling, and the pace of life in cities*. PNAS 104(17):7301–7306, 2007.
- [11] D. Helbing, A. Szolnoki, M. Perc, and G. Szabó. *Defector-accelerated cooperativeness and punishment in public goods games with mutations*. Phys. Rev. E 81:057104, 2010.
- [12] A. Szolnoki and M. Perc. *Phase diagrams for the spatial public goods game with pool punishment*. Phys. Rev. E 83:036101, 2011.
- [13] M. Ester *et al.* *A density-based algorithm for discovering clusters in large spatial databases with noise*. KDD’96, 1996.
- [14] L. Anselin. *Local indicators of spatial association — LISA*. Geographical Analysis 27(2):93–115, 1995.
- [15] A. Getis and J. K. Ord. *The analysis of spatial association by use of distance statistics*. Geographical Analysis 24:189–206, 1992.
- [16] M. Kulldorff. *A spatial scan statistic*. Communications in Statistics — Theory and Methods 26(6):1481–1496, 1997.
- [17] B. D. Ripley. *The second-order analysis of stationary point processes*. J. Applied Probability 13:255–266, 1976.

- [18] E. Goffman. *Asylums: Essays on the social situation of mental patients and other inmates*. Doubleday, 1961.
- [19] J. M. Robins, A. Rotnitzky, and L. P. Zhao. *Estimation of regression coefficients when some regressors are not always observed*. JASA 89(427): 846–866, 1994.
- [20] V. Chernozhukov *et al.* *Double/debiased machine learning for treatment and structural parameters*. Econometrics Journal 21(1):C1–C68, 2018.
- [21] S. Li. *Hotspot Policing in Toronto*. Capstone, 2022.
- [22] A. Hohl. *Adaptive Local Moran’s I for homicide cluster detection*. International Journal of Health Geographics, 2024.
- [23] V. S. Ruhela and Yoda. *clew: a content-addressable VCS with Rust, Zig, and C++ ports*. HADES-LLM internal paper, 2026.
- [24] V. S. Ruhela and Yoda. *MOIRAI: a function-per-file scientific-computing toolkit*. HADES-LLM internal paper, 2026.
- [25] V. S. Ruhela. *The Hawkes Paper: Markovian and non-Markovian self-exciting point processes for Toronto crime*. Private paper, 2026 ([/papers/hawkes-paper/](#)).
- [26] M. Wallace, J. Turner, C. Babyak, and A. Matarazzo. *Measuring Crime in Canada: Introducing the Crime Severity Index and Improvements to the Uniform Crime Reporting Survey*. Statistics Canada Catalogue 85-004-X, 2009.